

# Distributed File System

## Sharing

- Location-transparent sharing of data among workstations.
- A distributed file system subsumes fast and reliable underlying network connection, so that remote files appear local.
- Sharing of files is achieved by distinguishing two components:
  - **Client**
  - **Server**

# Distributed File System

## Motivation

### Goal

- To support efficient and transparent access to shared files maintaining data consistency.

### Problems

- Large number of mobile computers.
- Limited resources of mobile computers.
- Low and variable network bandwidth.
- Intermittent connection.
- Heterogeneity of h/w and s/w.
- Conventional RPC (ok for NFS) will not work.

### Solutions

- Replication of data (cloning, copying and caching)
- Data collection in anticipation of use (hoarding, prefetching)

# Distributed File System

## Consistency

### Problem of consistency

- In a loosely coupled distributed system all views on data will not be same.
- How view change should be propagated and to which users?
- Strong consistency (atomic update) out of question.

### Weak consistency

- Conflict resolution applied gradually afterwards to bring back consistency.
- Invalidation of data in mobile cache is problematic (disconnected mobiles).

# Distributed File System

## Update conflicts

### Conflict detection techniques

- Content independent: version numbering, timestamping
- Content dependent: dependency graph.

### Symmetry

- P2P relationship between client and server.
- Support in fixed n/w and/or mobile devices.
- One file or several file systems
- One or several namespaces for the files

# Distributed File System

## Connectedness

### Transparency

- Hides mobility support: applications do not require re-engineering.
- Users are oblivious of mobility.

### Consistency model

- Optimistic or pessimistic

### Caching and pre-fetching

- Single files, directories, subtree, partitions
- Permanent or only at certain points in time.

# Distributed File System

## Connectedness

### Data management

- Management of buffered data and copies.
- Requests for updates, validity of data.
- Detection of changes in data

### Conflict solution

- Application specific or general
- Errors.

# Distributed File System

## Storage systems for mobiles

- Storage systems: Coda, Bayou, and Rover
  - Bayou and Rover both support client mobility and weak network connectivity.
- Both employ replications, caching, and weak data consistency.
- **Bayou's focus:** application specific mechanism to detect conflict and resolve them.
- **Rover's focus:** both mobility-transparent and mobility-aware applications.
- **Coda's focus:** only application-transparent mobility.

# Distributed File System

## Goals of CODA

- High availability (n/w partitions, disconnected operations)
- Scalability
- Transparency
- Consistency
- Conflicting detection and resolution.
- Security



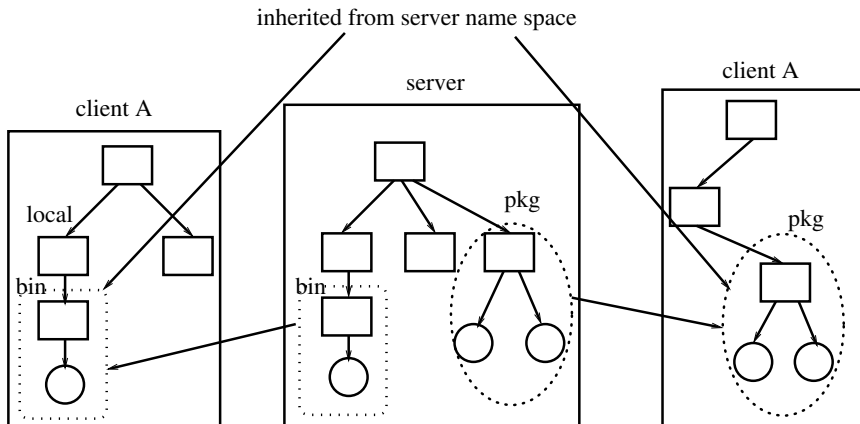
# Distributed File System

## Coda File System

- A client-side cache manager coupled with hoarding.
- Supports disconnected operation
- Makes sharing and collaborative computation easy in mobile computing environment.

# Overview of Coda

## Naming



# Overview of Coda

## Naming

- Files are grouped into volumes which is similar to unix partition but of smaller granularity.
- Basic units are volumes from which entire name space constructed.
- The construction takes place by mounting volumes at mount points.
- Mount points in Coda is a leaf node of one volume which becomes root node of another.

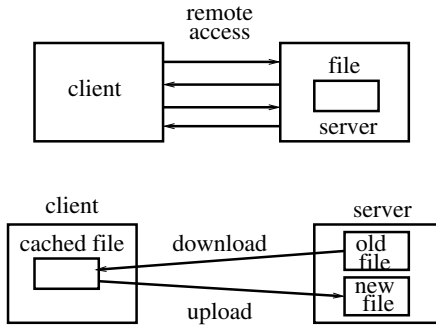
# Overview of Coda

## NFS and Coda Comparison

- A client uses mount commands on per server basis.
- File system should be exported.
- Based on RPC.
- Uses strong consistency.
- Knowing mount point of /coda suffices.
- Client-side component fetchs required file from appropriate server.
- Based on RPC2.
- Weak consistency.

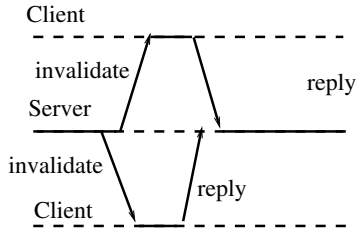
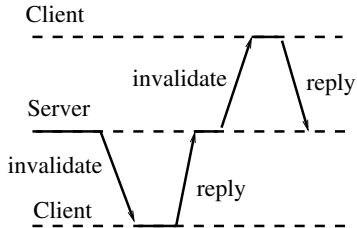
# Overview of Coda

## File access and update methods



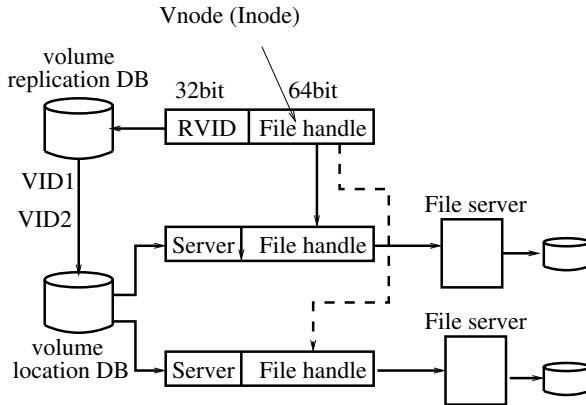
# Overview of Coda

## File access and update methods



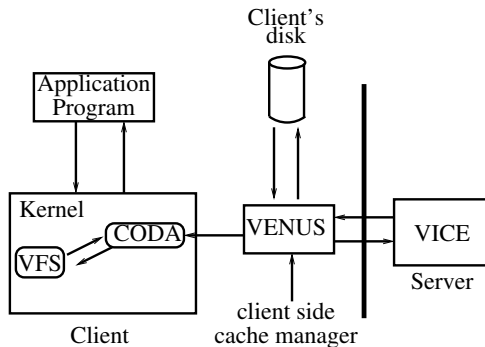
# Overview of Coda

## Fetching files from server



# Overview of Coda

## Components of Coda





# Overview of Coda

## Processing coda system call

- Any operation requiring access of shared objects is in granularity of files.
- User may be **connected** to network or operating in **disconnected** mode.
- Consider display of contents of a file in /coda directory.
- `cat <file>` generates a sequence of **system calls** with respect to the required files.
  - open, read, close.

# Overview of Coda

## Processing system calls

- The kernel at client requires VFS to open a file.
- VFS does the following:
  - Recognizes file argument is a coda resource.
  - Informs the kernel module for coda file system.
  - This module works basically a **in-kernel mini-cache**.

# Overview of Coda

## Role of in-kernel mini-cache

- Mini-cache is a cache of recently answered requests (for coda files) from VFS.
- If the file is locally available in cache area within the local coda file system, then VFS services the call.
- If file not locally available the mini-cache passes the request a user level program called **Venus**.
- Venus checks local disk, in the case of a cache miss contacts the server to locate the file.
- When file is located, Venus responds to kernel (mini-cache) which in turn services the client's request.

# Design of Coda

## Volume

- A client views coda as a single location independent Unix like file system.
- At the individual file server coda is mapped at the granularity of subtrees called **volumes**.
- A volume is a group of files and directories. located at one server
- The volume mappings are cached by Venus at the individual clients.

# Design of Coda

## VSG and AVSG

- Coda uses replications and callbacks for high availability and data consistency.
- The set of sites where replicas of file system volumes are placed called VSG.
- The subset of a VSG accessible to a client is called its AVSG.
- The overhead for maintaining replications at client is controlled by callback based cache coherence protocol.

# Design of Coda

## Replica

- Coda distinguishes between the replicas
  - **First class:** file copy at the server as
  - **second class:** cached at clients as
- If change occurs in server's replica then Venus notifies invalidation of cached copies.
- On server side updates are propagated to AVSGs then to missing VSGs.

# Design of Coda

## Disconnected operation

- When a client operates in disconnected mode: its AVSG is a null set.
- Clients can operate with the cached copies even when the AVSG is empty.
- Venus services the file system calls from its cache.
- The updates made by a client on the cached copies
- If cache miss occurs then failure reported.
- Update transmitted to server on reconnection.
- Venus recharges cache with server replication.

# Scalability

## Requirements

- The scalability requirements:
  - Ratio of clients to servers should be as large as possible.
  - Most of the load to be borne by clients.
  - Adding more clients does not significantly increase load on servers.
- OS running on a client filters out all the system calls that do not need interaction with servers.



# Scalability

## Filtering system calls

- Only open and close are forwarded to cache manager Venus.
- After a successful open, system calls read, write, seek can bypass Venus.
- Since servers are trusted and not loaded too much, they are responsible for preserving integrity and security of data.
- Most obvious approach to scale up is to employ caching at the clients.

# Replication and Availability

## Increasing availability

- Server replications increases availability of shared data.
- When a client is connected to at least one server it depends on the server replication.
- With many server replications, the period of disconnection gets minimized.
- Disconnected clients can switch to cached copies of shared data.
- But the difficulty is in preserving data consistency.

# Replication and Availability

## Consistency Versus Availability

- Maintaining consistency of cached replica at clients with the corresponding server replica which receive updates, should be done cleverly.
- So, a cache coherence protocol that combines consistency (first class with second class replicas) with scaling is required.
- Updates are generated by clients, so, consistency maintenance is complex.
- Several clients may use same replicas simultaneously and without each other's knowledge.

# Replication and Availability

## Restrictions

- Any rapid system-wide change is disallowed.
- Updates notified only if callback is registered. All the clients are not informed.
- There is no additional update overheads – like voting, quorum eliminated completely.
- So, consensus is sacrificed in preference to efficiency.

# Replication and Availability

## Restrictions

- The clients cache in granularity of entire files.
  - Files once opened successfully are assured to be available in local cache.
  - Subsequent operations on cached files are protected from network failures.
  - The idea of caching file in whole compatible with disconnected mode of operation.

# Replication and Availability

## Caching for availability

- Caching provides client autonomy to some extent.
- The files are cached in *whole*.
- Simplifies the failure model and makes disconnected mode of operation transparent to application.
- A cache miss can occur only at open.
- No cache miss for read, write, seek or close.

# Replication and Availability

## Caching for availability

- An application running disconnected client unaware of disconnection when the required file replicas are cached.
- Coda uses additional technique called *hoarding* to pre-cache files from the server.
- Hoarding is a large grain pre-fetching.

# Location Transparency

## Callbacks and cache coherency

- NFS does not provide location transparency.
- In AFS (on which coda is based), files are grouped into units or volumes that reside in a single server.
- Each volume contains mount points into other volumes on different servers.
- Mount point is a key (for look up in a distributed database) valid throughout an administrative domain called cell.
- So, with callback based cache coherence movement of files between servers could be possible.
- The client is required to know mapping of files to servers and to cache this dynamically changing map.



# Location Transparency

## Disconnection and failures

- Goal is to make client as much autonomous as possible, so that it becomes resilience to whole range failures.
- Complete autonomy means – personal computers operating in isolation.
- No sharing of data, antithetical to very purpose of distributed file system.
- Therefore, the coda goals were two-fold
  - using shared storage repository, and
  - continuance of client's operation even in the case of network failures.

# Replica Maintenance

## Replica types

- The server replicas represent the **first class** replicas.
- The client replicas represent the **second class** replicas.
- Second class replicas are known only to the client where they are located.
- Second class replicas increase overall data availability.
- Second class replicas must be periodically revalidated by synching.

# Replica Maintenance

## Replica control strategy

- The strategy may be either *pessimistic*, or *optimistic*.
- In pessimistic control strategy update is allowed only by one node in a unique n/w partition.
  - Prior to switching into disconnected mode, exclusive/shared lock on the cached object is to be acquired.
  - Lock relinquished only after reconnection.
  - Exclusive lock disallows read or write at any other replica.
  - Shared lock allows reads but disallows writes.

# Replica Maintenance

## Problems with locks

- Lock is restrictive, and infeasible in context of mobile clients as disconnection can be involuntary.
- Contention for lock can be a real problem even in voluntary disconnection.
- The duration of disconnection is unpredictable.
- Allowing exclusive lock on some objects only for a brief period of disconnection is acceptable.

# Replica Maintenance

## Lease lock

- Lease lock can solve problems in releasing locks by disconnected clients.
- But the purpose is defeated.
  - Disconnected client may be in the midst of update when it loses the lease.
  - Then all updates made by disconnected client have to be discarded.
  - It is also possible no other client is interested for the object when the disconnected client loses the lease for the lock.

# Replica Maintenance

## Optimistic control

- No exclusive lock is allowed, updates performed on cached objects.
- System has to be sophisticated to detect and resolve the update conflicts in reported client updates.
- Automatic conflict resolution is not only difficult but sometimes infeasible.
- It is possible to localize the conflicts and preserve evidences for a manual repair in extreme cases.

# Replica Maintenance

## Optimistic control

- The strategy provides a uniform model for every client when it operates in disconnected mode.
- The fact that the cached objects are in granularity of files also helps.
- In Unix like shared file system, usually the write access is with one user, other users share the file only for read accesses.
- The client with write access updates the cached file without any creating any update conflicts.
- The clients with read access may possibly get slightly outdated values, but applications can be tuned to handle such eventualities.

# Replica Maintenance

## Mixed control strategy

- Employ optimistic control strategy for disconnected mode and pessimistic strategy in connected mode.
- Requires users to be adaptive to the two conflicting approaches.
- Therefore, for transparent operation in disconnected mode, a uniform model of replica control is required.



# Replica Maintenance

## Visibility of updates

- In Unix, any modification to a file becomes immediately visible to all users.
- Preserving this semantics is incompatible with scalability and impractical when users operate in disconnected mode.
- So, coda settled for a diluted requirement of the visibility of updates under Unix file operation semantics.

# Replica Maintenance

## Visibility of updates

- The difference between a cached copy at a disconnected client from its latest copy depends on time difference between the last update of the copy and the current time.
- So, can be measured as a distance function of time difference.
- Revalidating cached file at the time of opening, and propagating updates at the time closing should be adequate for most applications.
- A successful open means the file being accessed is the latest replication available at the server from where the client fetched it.

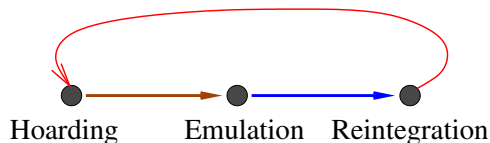
# Replica Maintenance

## Visibility of updates

- A server itself may have a stale copy.
- So, revalidation at the time of opening actually means fetching the latest of the replica at the accessible server site.
- Currency guarantee can be increased by querying the replica custodians on the fixed network when file opening is attempted at a client.
- This query is initiated by the server on which Venus docket the open call.

# Operation of Venus

## States of operation



- Venus fetches objects from accessible servers holding replicas of requested file objects.
- It propagates updates to the servers.
- Operates in three states

# Hoarding and Hoard Walk

## Factors Influencing Hoarding

- User's file reference behaviour in distant future.
  - Specifying reference profile can be helpful.
- Disconnection and reconnection are purely random in time.
- What is cost of cache miss during disconnection?
- Can cache size influences in balancing caching strategy?

# Hoarding and Hoard Walk

## Hoard Walk

- Dynamic nature of the above factors necessitates periodic re-evaluation of priorities.
- Limited cache size disturbs equilibrium as a result of normal activities.
  - Hoard walk should be performed periodically.
- Triggering hoard walk at the time of voluntary disconnection can be specified in profile.

# Emulation

## Venus works as sever's proxy

- Checks the access control.
- Generates *tfids* for newly created objects.
- Evaluates priorities of mutating objects.
- Reports or blocks a failures when a cache miss occurs.
- Records information to let the update activities done during disconnected operation to be *replayed* on reintegration.

# Emulation

## Cache management

- Caching is done using same priority algorithm as that for hoarding.
- Mutating objects directly update cache entries of involved objects.
- Cache entries of deleted objects are freed immediately.
- Cache entries of modified objects given infinite priorities.
- Cache miss by default reported unless explicitly masked.



# Emulation

## Replay log

- Per volume log is kept.
- All the operations during emulation state are logged.
- Operations replayed during the reintegration.
- Only updates are logged.
- Therefore, it is also known as change modify log (CML).
- Length of CML is reduced by applying optimization.

# Emulation

## Log optimization

- Operations followed by their inverses are purged from CML.
- A sequence of operations often replaced a single logical operation.
- For example, a `close` after an `open` installs a completely new file.
- `open`, the sequence of intervening `writes`, followed by a `close` replaced by a single `store`.
- `store` points to the cached copy of the file.
- A `store` invalidates the previous `stores` if any.
- So, recording the latest `store` is sufficient.

# Emulation

## Recovery guarantee

- In emulation state Venus must provides persistence guarantees:
  - A disconnected client to restart after a shutdown and continue from where it left off.
  - In the event of a crash while operating in disconnected state, the amount of data loss should not be more than that for an identical crash occurring in a connected state.

# Emulation

## RVM

- The *meta-data* used by Venus consists of
  - Cached directories and symbolic link contents,
  - Status blocks of cached objects of all types,
  - Replay logs, and
  - Hoard database (*hoard profile*) for the client.
- Meta data are mapped into address space of Venus as the RVM.
- The contents of cached files are stored as regular Unix files on client's local disk.

# Emulation

## RVM

- Transactional operation guarantees consistency.
- Crash recovery is ensured by log of the operation.
- Only non-nested local transactions possible.
- To reduce commit latency commits are labelled as `no-flush`
- To ensure persistence of `no-flush` commits log is periodically flushed.
- So RVM provides bounded persistence
  - Where bound is the period of between two flushes.

# Emulation

## Resource exhaustion

- Non-volatile storage may exhaust during emulation.
  - File cache become full with modified files.
  - RVM space for log becomes full.
- In case of file cache space is freed by deleting or truncating modified files.
- When log space is full no mutation is allowed until reintegration.
- Non mutating operations are allowed.

# Reintegration

## Role of Venus

- It changes from pseudo server state to cache manager.
- Changes during emulation state is propagated.
- Cache is updated.
- Reintegration is performed on a volume at a time.
- All update activities on volume is suspended until completion.

# Reintegration

## Replay algorithm

- First step is to allocate permanent fids.
  - Can be avoided if a small supply of fids cached in advance during hoarding state.
- In second step replay log is transferred in parallel to AVSG and execute independently.
- Every server will execute replay log in one transaction.



# Reintegration

## Replay algorithm

- Consists of 4 phases: log parsing, validation, backfetching and commit
- Each operation is validate and executed.
  - Consists of integrity, conflict detection, protection and disk space checks.
  - Except for store all other operations are executed directly.
- For store a shadow file is created and meta data updates done.
  - Data transfer deferred to phase 3.

# Reintegration

## Replay algorithm

- Third phase is backfetching: for data transfer.
- Final phase is releasing all locks.
- If reintegration succeeds replay log is freed.
- If reintegration fails replay log is written to a local replay file and cache entries are purged.
- So subsequent references fetches current contents of AVSG.
- Manual tool is provided for inspection of replay file.
- Selective replay is possible.